Beyond Forecast Accuracy: Comparing Probabilistic Demand Forecasts by the Inventory They Require across Retail Settings and Periods

Bùi Anh Duy, Khổng Trần Gia Huy, Nguyễn Thị Cẩm Hương, Trương Long

FPT University, Ho Chi Minh City, Vietnam · anhduy290605@gmail.com; huyktgse180313@fpt.edu.vn; huongntc2@fpt.edu.vn; longt5@fpt.edu.vn

ABSTRACT

The forecast with the lowest error is not necessarily the forecast that supports the lowest inventory. We test this distinction using seven probabilistic methods—three statistical approaches and four gradient-boosting variants—in a weekly order-up-to simulation with optional liquidation. The experiments retain all eligible series from M5 (physical stores) and VN1 (multi-vendor e-commerce) and cover three non-overlapping 26-week windows per dataset. Because observed inventory costs are unavailable, we compare stock requirements at equal realised fill rates and report bootstrap confidence intervals. On M5, LightGBM quantile regression (LGB-Q) yields the lowest volume-weighted quantile loss. At a fill rate of 0.94, it uses 3–11% less inventory than the safety-stock, conformal, TSB-NB and exponential-smoothing alternatives in each window. The M5 result also persists across the tested lead times and LightGBM settings. VN1 gives a different ranking in different windows: TSB-NB uses 13–16% less inventory for stable-demand series, but fails to reach a fill rate of 0.92 on growing-demand series in two windows. Historical ADI–CV² classes and the tested clusters do not consistently identify which method to use. Averaging the LGB-Q and TSB-NB quantiles, a post hoc choice, reduces VN1 inventory by 6–11% relative to LGB-Q at a fill rate of 0.92 and remains within 1.5% on M5; a median of all seven methods, specified before testing, saves 7–8% on VN1 and stays within 1.6% on M5. Because these combinations lack an independent test period, our main evidence concerns the single methods. The N-HiTS baseline leads none of the inventory comparisons. Quantile-based liquidation changes fill rate by at most 0.02 percentage points, but has no consistent economic advantage over 26 weeks under the assumed costs.

CCS CONCEPTS • Applied computing → Forecasting; Supply chain management • Computing methodologies → Machine learning; Cluster analysis

Additional Keywords and Phrases: probabilistic forecasting; inventory management; intermittent demand; quantile regression; demand segmentation; liquidation; M5; VN1

ACM Reference Format: Bùi Anh Duy, Khổng Trần Gia Huy, Nguyễn Thị Cẩm Hương, and Trương Long. 2027. Beyond Forecast Accuracy: Comparing Probabilistic Demand Forecasts by the Inventory They Require across Retail Settings and Periods. In Proceedings of the 12th International Conference on Intelligent Information Technology (ICIIT 2027). ACM, New York, NY, USA.

1 INTRODUCTION

An order based on the expected value of demand cannot describe the risk of a stockout. That risk matters: insufficient inventory loses sales, while excess stock incurs holding costs and may later be liquidated. Forecast quantiles allow an order-up-to level to be set for a chosen service target [3]. Whether the quantiles are statistically accurate, however, is a separate question from how much stock their use requires.

The difference matters when selecting a forecasting method. M5 ranks forecasts using weighted RMSSE and weighted scaled pinball loss [1]; its uncertainty track was not designed around an inventory decision [3]. Work with VN1 has likewise focused on accuracy and computation [4]. Inventory simulations can yield a different ranking [5, 6, 23]. In particular, Theodorou et al. [7] report discrepancies on M5 when lost sales are costly, especially for intermittent series. They identify online retail as a setting in need of further testing.

Three features of the existing evidence motivate our design. First, M5 inventory evaluations have used selected subsets: 8,000 high-volume series [8], volatile food products [9], or series with at least 2.5 years of history [7]. These restrictions leave some sparse series outside the analysis. Second, most comparisons use a single dataset and test period [5, 7–9, 23], so period-to-period variation is poorly characterised. Third, public sales panels do not provide historical inventory positions or unit costs. Any simulated monetary outcome must consequently depend on assumptions about those missing quantities [7].

We hold the inventory policy fixed and vary its demand forecasts rather than introduce another forecasting architecture. Seven methods are evaluated on every eligible series in M5 and VN1, with weekly decisions simulated over three 26-week windows in each dataset. Inventory is compared at matched realised fill rates; liquidation is assessed through the salvage price needed to break even under stated cost assumptions. Additional experiments vary LightGBM hyperparameters and replenishment lead times and include a neural baseline. We also test whether pre-test demand characteristics can support method selection. Finally, we report an exploratory average of two forecast distributions, chosen after seeing the initial results, and a pre-specified combination and ablation of all seven methods.

2 RELATED WORK

Retail forecasting benchmarks. In M5's accuracy track, all 50 highest-ranked teams used LightGBM, and the winning entry combined models fitted with Tweedie loss [2]. The leading uncertainty-track entry trained a separate LightGBM model for each quantile [3]. We therefore include both point-based and directly estimated quantiles. Zanotti [4] compared global-model ensembles on M5 and VN1 in terms of forecast accuracy and computing cost, without simulating stock decisions.

Intermittent demand. TSB updates the probability of non-zero demand from one period to the next, allowing the forecast to decline when sales stop [10]. Average inter-demand interval (ADI) and squared coefficient of variation of positive demand sizes (CV^2) are conventional descriptions of demand intermittency [11]. For M5 we use class boundaries ADI = 4/3 and $CV^2 = 0.5$ [1]. Damato et al. [12] found no established global-model architecture that dominates on intermittent data.

Forecasts used for inventory. Goltsos et al. [6] document the frequent separation of forecasting studies from the replenishment decisions they inform. Kourentzes et al. [13] fit exponential smoothing with an inventory-based criterion; van der Haar et al. [14] estimate order quantities directly. On M3 series, Petropoulos et al. [24] found forecast combinations effective in a rolling order-up-to simulation. Evidence from M5 is mixed: simple empirical quantiles perform well in the 28-day daily-review experiment of Spiliotis et al. [23]; Wang et al. [5] study forecast combinations with a single-period newsvendor; Zabraoui et al. [9] apply reinforcement-learning stock policies to selected products; and Mohammed et al. [8] evaluate safety stock for high-volume series with an illustrative cost calculation.

Theodorou et al. [7] are the closest point of comparison. They simulate lost-sales inventory control for one year on 7,783 daily M5 product–state series. Their design covers 13 forecasting methods and 24 service-target, review-period and lead-time combinations. Quantiles are constructed from point forecasts and past forecast errors. Although error and inventory trade-off curves generally align, inventory-cost rankings can favour less accurate methods when stockouts carry a high penalty.

We retain eligible series across demand classes, use two retail settings and three test windows, and include a liquidation decision. We also estimate quantiles directly rather than deriving all of them from residuals. Inventory is compared at the same realised fill rate, because a common nominal quantile level does not give all methods the same achieved service [7]. Neither multi-period stock simulation nor decision-based forecast evaluation is new [5, 7, 9, 14]; the contribution here is the comparison across datasets, periods and decision settings.

3 METHOD

3.1 Data

M5 [1] reports daily unit sales for 3,049 products in 10 Walmart stores, giving 30,490 product–store series. Aggregating complete weeks yields 277 weekly observations. VN1 [21] contains 15,053 client–warehouse–product series spanning 196 weeks of multi-vendor e-commerce sales; despite its name, the dataset is not from Vietnam. Eligibility is determined separately for each test window: a series must start at least 13 weeks before that window and record at least one earlier sale.

Table 1. Dataset characteristics measured before the main test window.

Property	M5	VN1
Retail setting	One retailer; 10 physical stores	Multi-vendor e-commerce; 46 vendors; 328 warehouses
Total series	30,490	15,053
Eligible series (main / second / third)	30,381 / 29,917 / 28,824	13,844 / 11,442 / 9,383
Number of weeks	277	196
Zero-demand weeks since series start	24.4%	49.1%
Smooth / erratic / intermittent / lumpy	53.0 / 5.8 / 33.9 / 7.3%	18.8 / 13.7 / 47.2 / 20.3%
Additional variables	Prices, events, SNAP days, product hierarchy	Prices in weeks with sales only

A single processing pipeline handles both panels. The source data differ in the definition of a series start, but the resulting dates agree for 99.5% of M5 series. Their demand characteristics are nonetheless distinct: a classifier using nine demand features identifies the dataset with an AUC of 0.91 (Section 5.7).

3.2 Forecast targets and methods

Each week, we predict quantiles $q \in \{0.50, 0.80, 0.90, 0.95, 0.99\}$ of demand accumulated over horizon h. Replenishment uses h = 3 weeks and liquidation h = 13 weeks. Quantiles are fitted to cumulative demand rather than obtained by summing separate weekly quantiles, which would generally be incorrect.

EMP: empirical quantiles of rolling h-week demand totals observed during the preceding 104 weeks.

ETS: simple exponential smoothing and a normal approximation for the h-week demand total [18].

TSB-NB: TSB [10] with a negative-binomial predictive distribution [25], matched to the estimated mean and variance of h-week demand.

LGB-T: a global LightGBM point model fitted with Tweedie loss [15], with a normal safety-stock adjustment derived from validation residuals for each series.

LGB-C: the same point model, with split-conformal quantiles [17] based on residuals pooled within demand classes.

HGB-Q: histogram gradient boosting [19] fitted with quantile loss to 300,000 sampled training rows. It provides a second direct-quantile implementation.

LGB-Q: LightGBM quantile regression [16], fitting a separate global model for each quantile.

An additional combination, C2, specified post hoc after inspecting the results, averages the corresponding LGB-Q and TSB-NB quantiles (Section 5.5). All four global gradient-boosting methods have the same input categories: sales lags, rolling mean and variation, proportion of zero-sales weeks, time since last sale, calendar terms, price and series scale. Targets are scaled to the series level. M5 additionally supplies events, SNAP days and product hierarchy. No hyperparameter search is performed. LGB-Q, LGB-T and LGB-C share one LightGBM configuration: learning rate 0.05, 63 leaves with no depth limit, at least 200 samples per leaf, L2 penalty 1, feature and row subsampling fractions of 0.8, and at most 1,000 boosting rounds with early stopping after 50 rounds. Each refit trains on up to 104 weekly origins and validates on the following 13; LGB-T and LGB-C use Tweedie loss with variance power 1.1. HGB-Q uses learning rate 0.1, at most 300 iterations, 63 leaves and at least 200 samples per leaf. Sensitivity checks appear in Section 5.8.

3.3 Inventory decision policy

We simulate weekly review with a separate lost-sales inventory process for each series. At the start of a week, orders from L weeks earlier arrive. Inventory position IP includes on-hand and outstanding units. If liquidation applies, stock above threshold T is removed. The replenishment order is $\max(S - IP, 0)$, where S is the τ-quantile of demand over L + 1 weeks; no new order is placed in a week with liquidation. Demand then occurs, and any unmet units are recorded as lost sales. Unless otherwise stated, L = 2 weeks.

Four policies share the replenishment forecasts: no liquidation; quantile liquidation, with T set to the method's 0.95-quantile of 13-week demand; a fixed threshold of 26 times mean weekly sales in the preceding 26 weeks; and dead13, which removes all stock after 13 consecutive weeks without sales and suspends ordering until a new sale is recorded.

3.4 Evaluation metrics

Forecast accuracy. For a series, scaled quantile loss divides pinball loss by the in-sample naïve forecast error. A small denominator makes its cross-series mean unstable on VN1 (Section 5.1). We therefore report volume-weighted quantile loss (wQL: total pinball loss divided by total demand), the median series-level scaled loss, and the mean rank of each method across series. Rank differences are evaluated with the Friedman test and Nemenyi critical difference [20].

Inventory. Each 26-week window starts with a four-week warm-up, leaving 22 weeks for evaluation. Fill rate is the fraction of demanded units sold; inventory is measured in weeks of demand. We vary the order-up-to quantile level τ over the five reported values, then interpolate along each inventory–fill-rate curve to find the stock required at a specified realised fill rate [7, 25]. A method's frontier rank is its mean inventory rank across fill targets of 0.90–0.98. The 200 series-level bootstrap resamples reflect uncertainty across series, not uncertainty about performance in future periods.

Liquidation. For each policy we simulate each series twice from the same starting stock, with and without liquidation, and write Δ for the difference between the two runs over the test window. With unit cost normalised to 1, ΔO is the change in units ordered, ΔS the change in units sold, ΔI the change in inventory (unit-weeks), ΔE the change in the end-of-horizon position (on-hand plus on-order units, valued at cost) and Q the number of units liquidated. With weekly holding rate h (annual rate / 52) and gross margin m (selling price (1 + m) times cost), liquidating at salvage ratio s is at least as profitable as holding when $s \cdot Q + (1 + m) \cdot \Delta S - \Delta O - h \cdot \Delta I + \Delta E \geq 0$. The break-even salvage ratio is therefore $s^* = [\Delta O - \Delta E + h \cdot \Delta I - (1 + m) \cdot \Delta S] / Q$. A value $s^* < 1$ means that liquidating below cost still pays off; $s^* > 1$ means that liquidation pays off only if stock can be sold above cost. Valuing the end position at cost favours holding, so the reported

s* is an upper bound. Calculations cover annual holding costs of 10–40% and gross margins of 30–100%. Neither actual salvage prices nor inventory costs are observed in these datasets.

4 EXPERIMENTAL SETUP

For each dataset, the main window comprises the final 26 weeks; two earlier windows are formed by removing the last 26 and 52 weeks, respectively. Within each window, the models are refitted at the beginning of two 13-week blocks, while forecasts and inventory decisions are generated weekly. Training is repeated for all three windows. Experiments ran on an Intel Core i5-12450H laptop CPU with 16 GB of RAM. Code and result files: https://github.com/anhduyBA/forecast-to-decision-benchmark.

5 RESULTS

5.1 Forecast accuracy

Table 2. Three-week forecast accuracy, by dataset and window. Lower values indicate better performance; wQL is volume-weighted quantile loss.

Method	Window	M5 wQL	M5 median	M5 rank	VN1 wQL	VN1 median	VN1 rank
EMP	main	0.115	0.174	4.85	0.172	0.202	4.72
EMP	second	0.120	0.176	4.88	0.146	0.207	4.91
EMP	third	0.126	0.176	4.95	0.197	0.239	4.90
TSB-NB	main	0.088	0.155	4.22	0.129	0.152	3.25
TSB-NB	second	0.091	0.150	3.99	0.088	0.150	3.30
TSB-NB	third	0.098	0.150	4.12	0.137	0.186	3.31
ETS	main	0.084	0.153	4.12	0.129	0.174	4.50
ETS	second	0.089	0.152	4.09	0.092	0.173	4.54
ETS	third	0.095	0.150	4.12	0.136	0.198	4.28
LGB-T	main	0.083	0.152	3.87	0.121	0.188	4.68
LGB-T	second	0.088	0.151	3.90	0.086	0.166	4.45
LGB-T	third	0.094	0.148	3.89	0.126	0.213	4.51
LGB-C	main	0.086	0.159	4.50	0.120	0.184	4.43
LGB-C	second	0.091	0.160	4.60	0.091	0.173	4.48
LGB-C	third	0.096	0.155	4.48	0.125	0.231	4.65
HGB-Q	main	0.077	0.146	3.35	0.104	0.153	3.23
HGB-Q	second	0.081	0.145	3.38	0.079	0.140	3.19
HGB-Q	third	0.087	0.142	3.29	0.104	0.184	3.20
LGB-Q	main	0.076	0.145	3.10	0.103	0.152	3.18
LGB-Q	second	0.080	0.144	3.17	0.078	0.140	3.13
LGB-Q	third	0.086	0.142	3.15	0.102	0.184	3.14

Nemenyi critical difference ($\alpha = 0.05$), main / second / third: M5 0.052 / 0.052 / 0.053; VN1 0.077 / 0.084 / 0.093. At h = 13, TSB-NB has the best mean VN1 rank in every window (2.92 / 2.99 / 3.37).

LGB-Q has the lowest wQL at h = 3 in both datasets and all three windows (Table 2); HGB-Q is within 0.002. The same wQL ordering holds within each demand class. On M5, the average per-series rank also favours LGB-Q in every window, by more than the Nemenyi critical difference relative to HGB-Q.

The VN1 rankings are less clear at series level. At h = 3, LGB-Q, HGB-Q and TSB-NB are separated by no more than 0.17 mean ranks, with no significant LGB-Q–HGB-Q difference under the Nemenyi criterion. At h = 13, TSB-NB ranks first in all three windows. LGB-Q's aggregate wQL advantage thus does not extend uniformly to individual VN1 series. Its mean scaled loss on VN1 is 0.336, 0.539 and 2.008 across the windows, inflated by a few series with very small naïve-error denominators; wQL remains 0.103, 0.078 and 0.102.

5.2 Inventory at equal realised fill rate

M5. LGB-Q ranks first on the inventory frontier in all bootstrap resamples across the three windows. At fill rate 0.94, it needs 3.3–10.6% less stock than LGB-T, LGB-C, TSB-NB and ETS, with the reported intervals excluding zero (Table 3). HGB-Q's inventory requirement differs by 0.2–1.4%.

VN1. The preferred method depends on the window. LGB-Q ranks first in the main window, although its stock advantage over LGB-T, LGB-C and TSB-NB at fill rate 0.94 is not established by the reported confidence intervals. In the second window, TSB-NB ranks first; TSB-NB, LGB-T and LGB-C need 4.1–8.8% less inventory than LGB-Q. In the third, LGB-Q again leads, while TSB-NB, ETS and LGB-T do not attain fill rate 0.94 at any evaluated τ. Lowest wQL does not always mean lowest stock on VN1.

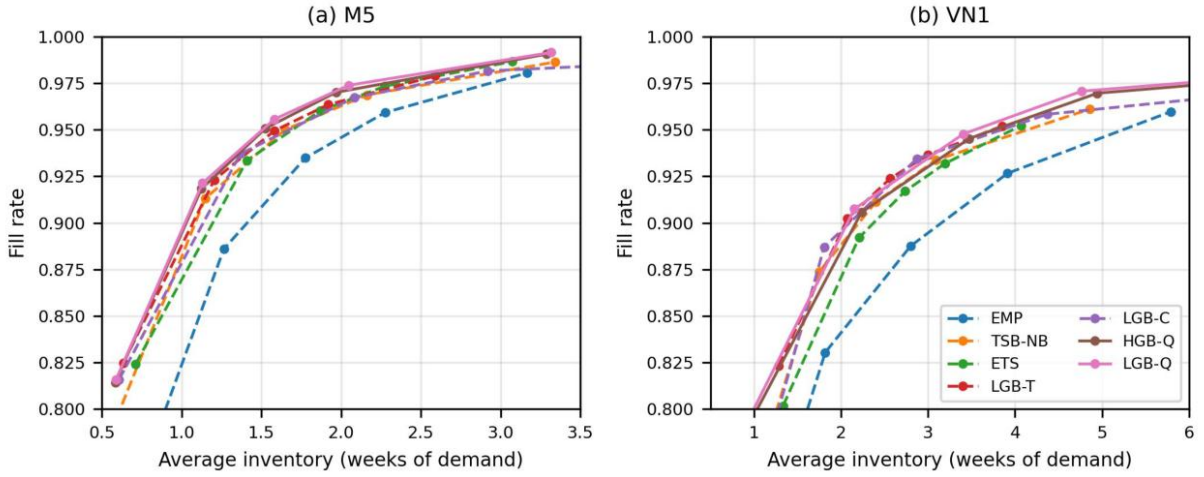


Figure 1. Inventory in weeks of demand against realised fill rate in the main window: (a) M5, (b) VN1. τ ranges from 0.50 to 0.99; solid curves represent direct quantile methods.

Table 3. Inventory needed for fill rate 0.94. Δ = I_method / I_LGB-Q − 1 (%); positive values favour LGB-Q. Brackets show 95% bootstrap intervals; "—" means the method did not reach 0.94.

Set	Window	LGB-Q Inv.	Frontier rank	P(best)	HGB-Q Δ %	LGB-T Δ %	LGB-C Δ %	TSB-NB Δ %	ETS Δ %
M5	main	1.38 [1.34, 1.41]	1.00	1.00	+1.2 [+0.9,+1.6]	+5.1 [+4.4,+5.9]	+4.5 [+3.2,+5.8]	+10.0 [+9.1,+11.0]	+10.6 [+9.8,+11.5]
M5	second	1.49 [1.45, 1.53]	1.20	1.00	+1.4 [+1.2,+1.7]	+6.8 [+6.2,+7.6]	+8.2 [+6.8,+9.3]	+7.5 [+6.7,+8.3]	+8.0 [+7.1,+8.9]
M5	third	1.74 [1.69, 1.78]	1.00	1.00	+0.2 [−0.2,+0.5]	+3.3 [+2.7,+4.0]	+7.5 [+6.3,+8.4]	+4.9 [+4.2,+5.7]	+5.9 [+5.0,+6.6]
VN1	main	3.17 [2.78, 3.57]	1.40	0.93	+4.5 [+2.9,+6.7]	+0.9 [−7.5,+9.9]	+1.7 [−6.3,+8.7]	+10.5 [−0.5,+19.8]	+12.0 [+3.8,+19.7]
VN1	second	1.79 [1.52, 2.05]	3.80	0.00	−1.2 [−2.5,+0.2]	−7.7 [−11.8,−3.8]	−4.1 [−7.3,−0.9]	−8.8 [−12.8,−4.9]	−0.4 [−5.7,+4.2]
VN1	third	3.31 [2.72, 3.85]	1.25	0.95	+1.3 [−0.7,+3.5]	—	+11.5 [+1.6,+116.5]	—	—

LGB-Q requires the least inventory in the smooth, intermittent and lumpy M5 classes in all three windows. On VN1 it leads the lumpy class consistently, but not the smooth class: TSB-NB leads that class twice and LGB-C once.

5.3 Period variation on VN1

Table 4. Demand groups defined retrospectively from test-period demand relative to the preceding 26 weeks (stable: 0.67–1.50×; growing: >1.50×). Δ compares TSB-NB inventory with LGB-Q at fill rate 0.92.

Dataset	Window	Growing share of test demand	Stable Δ %	Growing Δ %	All series Δ %
M5	main	20.7%	+8.5	+12.5	+10.4
M5	second	28.0%	+7.7	+7.7	+5.5
M5	third	19.4%	+3.9	+5.7	+6.1
VN1	main	31.5%	−14.4	TSB-NB not reached	+5.0
VN1	second	22.8%	−13.3	−10.5	−11.3
VN1	third	40.4%	−15.6	TSB-NB not reached	+25.1

For this analysis only, we group series according to the realised ratio of test-window demand to demand in the preceding 26 weeks. Ratios of 0.67–1.50 define the stable group; ratios above 1.50 define the growing group. Because grouping uses outcomes from the test window itself, it cannot be applied as a forecasting-time selection rule.

The stable-demand VN1 group favours TSB-NB in each window: at fill rate 0.92, it uses 13.3–15.6% less inventory than LGB-Q (Table 4). For growing VN1 series, TSB-NB fails to reach that fill rate in the main and third windows, while LGB-Q reaches it. Growing series account for 31.5% and 40.4% of test demand in those windows, compared with 22.8% in the second window, where TSB-NB leads overall. This pattern associates the ranking with demand composition but does not establish causation. On M5, LGB-Q needs less stock in both groups throughout.

5.4 Prospective method selection using demand groups

Table 5. Selection rules trained on one window and evaluated on the next. Numbers are inventory changes (%) relative to always using LGB-Q at fill rate 0.94; negatives favour the rule. Brackets give 95% bootstrap intervals.

Selection rule	M5 third → second	M5 second → main	VN1 third → second	VN1 second → main
Always TSB-NB	+7.5 [+6.7,+8.6]	+10.0 [+8.9,+10.9]	−8.8 [−12.6,−5.1]	+10.5 [+1.9,+20.4]
Best single method in previous window	+0.0 [+0.0,+0.0]	+0.0 [+0.0,+0.0]	+0.0 [+0.0,+0.0]	+10.5 [+1.9,+20.4]
By ADI–CV² class	+0.1 [−0.0,+0.2]	+0.0 [+0.0,+0.0]	−5.6 [−8.1,−3.9]	+5.1 [−0.3,+10.2]
By segment	+0.1 [−0.1,+0.3]	+0.0 [+0.0,+0.0]	−0.1 [−0.7,+0.5]	+10.5 [+1.9,+20.4]
By class × segment	+0.1 [−0.0,+0.3]	+0.0 [+0.0,+0.0]	−5.6 [−8.1,−3.8]	+4.8 [−0.1,+12.0]
Oracle by segment (not deployable)	+0.0 [+0.0,+0.0]	+0.0 [+0.0,+0.0]	−8.9 [−12.8,−5.2]	+0.0 [+0.0,+0.0]

We next ask whether information available before a window can identify which forecaster to use. In addition to ADI–CV² classes, we cluster series using eight historical features: level, order size, ADI, CV², time since sale, trend, recent relative level and peak size. ADI, CV² and time since sale are log-transformed; all eight features are then centred on their medians, divided by their interquartile ranges and clipped at ±4, with these statistics taken from the fitting window. The clustering fitted on the earliest window is carried forward unchanged. Predefined criteria select a three-component Gaussian mixture (regular, fading, sparse) that is stable across refits (adjusted Rand index 0.94) but overlapping (silhouette 0.21).

Neither grouping gives reliable improvements (Table 5). On M5, choosing the best method for each group in the previous window changes inventory by no more than 0.3% relative to LGB-Q. For VN1, class-based selection first saves 5.6% and then requires 5.1% more stock; segmentation produces changes of −0.1% and +10.5%. Even a retrospective oracle choosing the best method within each segment cannot beat the best single method of that window. Growing series appear across the groups rather than in a distinct cluster: they represent 18–31% of each VN1 segment and 14–39% of each ADI–CV² class.

5.5 Inventory performance of averaged quantiles

The complementary VN1 results for LGB-Q and TSB-NB suggest testing their average, C2, the arithmetic mean of the two methods' quantiles. We chose these two inputs after inspecting their relative performance, so C2 is exploratory and post hoc. To test combination without that choice, we fixed a second design before running it: the mean (E7) and the median (E7-med) of all seven methods' quantiles, leave-one-out ablations of E7, and a prospective rule that applies the lowest-regret mean subset of one window to the next.

At fill rate 0.92, C2 uses 5.9%, 10.5% and 7.1% less VN1 inventory than LGB-Q in the main, second and third windows and E7-med 7.6%, 8.2% and 7.1% less, with all bootstrap intervals excluding zero. On M5 both stay within 1.6% of LGB-Q, whereas E7 needs up to 3.1% more and so misses our preset criterion. Every combination has a higher wQL than LGB-Q. Removing LGB-Q or HGB-Q from E7 raises its stock by 1.0–4.5%, removing ETS lowers it in all six runs, and the lowest-regret subset of every run contains LGB-Q and TSB-NB. The prospective rule saves 7.8–8.2% on VN1 in both transitions but needs 0.9% more on M5 in one. C2 also saves VN1 inventory at one- and four-week lead times (Section 5.8).

5.6 Liquidation outcomes

Table 6. Changes relative to no liquidation when LGB-Q supplies the forecasts. Break-even salvage ratios s are calculated under the specified holding-cost and margin assumptions.*

Set	Policy	Window	Liquidated (% demand)	Inventory change (%)	Fill-rate change (pp)	Stockout change (%)	s* range
M5	Quantile	main	0.02	−0.1	−0.00	+0	Unstable (few units)
M5	Quantile	second	0.01	−0.1	−0.00	+0	Unstable (few

Set	Policy	Window	Liquidated (% demand)	Inventory change (%)	Fill-rate change (pp)	Stockout change (%)	s* range
							units)
M5	Quantile	third	0.02	−0.2	−0.00	+0	Unstable (few units)
M5	Fixed	main	0.79	−2.3	−0.29	+7	1.06–1.32
M5	Fixed	second	1.08	−3.0	−0.20	+5	1.01–1.18
M5	Fixed	third	1.21	−2.7	−0.29	+7	1.03–1.23
M5	Dead13	main	0.67	−2.5	−0.65	+13	1.35–2.08
M5	Dead13	second	0.82	−4.9	−0.46	+10	1.21–1.66
M5	Dead13	third	0.90	−3.3	−0.54	+13	1.19–1.70
VN1	Quantile	main	1.18	−4.3	−0.02	+0	0.91–1.04
VN1	Quantile	second	0.27	−1.7	−0.02	+1	0.92–1.21
VN1	Quantile	third	1.74	−7.4	−0.02	+0	0.91–1.04
VN1	Fixed	main	2.76	−6.1	−0.67	+20	1.00–1.22
VN1	Fixed	second	1.47	−5.0	−0.11	+18	0.97–1.07
VN1	Fixed	third	1.91	−4.4	−0.62	+17	1.03–1.29
VN1	Dead13	main	1.21	−4.0	−0.09	+23	1.03–1.37
VN1	Dead13	second	0.76	−3.8	−0.08	+22	1.06–1.32
VN1	Dead13	third	0.47	−2.3	−0.07	+18	1.16–1.60

With LGB-Q forecasts, the quantile rule lowers fill rate by no more than 0.02 percentage points in any window (Table 6). It rarely removes stock on M5; on VN1, it reduces inventory by 1.7–7.4%. Fixed and dead13 liquidation have larger service effects, increasing stockout rates by 5–20% and 10–23%, respectively.

None of the policies has a consistent economic advantage over the 26-week simulation under the assumed costs. For quantile liquidation on VN1, the break-even salvage ratio ranges from 0.91 to 1.21 times unit purchase cost. Ratios for the remaining policies have a lower bound of 0.97. These figures are upper-bound thresholds from the cost model of Section 3.4, not measured resale revenues or realised profits. In practice the rule is direct: liquidate only if the expected recovery ratio r on the secondary market (salvage price divided by unit cost) exceeds s*. With the VN1 quantile-rule thresholds of 0.91–1.21, a channel that recovers 60% of cost would not justify liquidation within 26 weeks, whereas one that recovers about unit cost could. Obsolescence, scarce storage space or a longer horizon would lower s* and are not modelled here.

5.7 Sensitivity to dataset differences

We tested whether specific differences between M5 and VN1 account for the contrasting rankings. Reweighting VN1 to approximate M5's demand-feature distribution reduces the mean standardised difference from 0.32–0.39 to 0.02–0.03 without altering the VN1 findings. Removing price variables from VN1, where prices are available in fewer than half of the weeks, or limiting M5 to its 21 features shared with VN1 changes LGB-Q inventory by less than 1% in every window. The tested covariate and preprocessing differences do not account for the results; other dataset differences remain untested. Further implementation details are in the repository.

5.8 Robustness checks and a neural baseline

Hyperparameters. For h = 3 in the main window, we refitted LGB-Q and LGB-T with small (15 leaves), large (255 leaves) and regularised (31 leaves, minimum 500 samples per leaf, L2 penalty 10) trees. On M5, LGB-Q needs 4.7–6.8% less inventory than LGB-T and 7.5–11.0% less than TSB-NB at fill rate 0.92 under all four configurations, and its own inventory changes by less than 3%. The VN1 ranking persists for three settings; with small trees, LGB-Q's advantage over TSB-NB is no longer resolved and LGB-T uses 8.9% less inventory than LGB-Q.

Lead time. We repeated the equal-fill comparisons at lead times of one and four weeks, corresponding to forecast horizons of two and five weeks, in every test window. HGB-Q was omitted from this check. On M5, across all nine window–lead-time combinations, LGB-Q requires less inventory than LGB-T (2.5–7.4%), LGB-C (4.5–10.9%), TSB-NB (4.9–14.8%) and ETS (5.9–13.4%) at fill rate 0.94. Its regret remains at or below 0.2%. For VN1, the post hoc C2 uses 3.2–10.9% less inventory than LGB-Q at fill rate 0.92 in all nine combinations, with confidence intervals excluding zero in seven. Across the 18 dataset–window–lead-time runs, C2's maximum regret relative to the best of the six single methods is 2.0%; LGB-Q's is 12.6%, in the second VN1 window at four weeks of lead time.

Neural baseline. We evaluated a univariate N-HiTS network [26] separately from the seven-method benchmark. It predicts quantiles of cumulative h-week demand with a multi-quantile loss, is refitted every 13 weeks and uses fixed

hyperparameters. Its wQL lies between LGB-Q and TSB-NB in all six runs, and its 0.99-quantile coverage is 0.986–0.992. It never leads on inventory. At fill rates of 0.90–0.94 on M5 it needs 2.6–10.0% more stock than LGB-Q; its frontier ranks on VN1, among eight single methods, range from 2.8 to 4.4. Including N-HiTS leaves the post hoc C2's maximum regret at 0.9%. Because we did not test architectures such as DeepAR or PatchTST, this CPU-budgeted baseline is not evidence about deep-learning methods as a class.

Computational cost. All methods except Chronos-2 forecast one 26-week window in minutes on the laptop CPU: about 1–2 minutes for TSB-NB, 10–21 minutes for LGB-Q on M5 (five quantile models per refit) and about 7 minutes for N-HiTS. Zero-shot Chronos-2 [22] needs no training, but its CPU inference took 5.4 hours per M5 window and 32–45 minutes per VN1 window, while needing 14.7% more inventory than LGB-Q on M5 and 1.2–48.1% more on VN1 at fill rate 0.92. For weekly review, computation does not constrain the statistical and boosting methods, and the costlier zero-shot model did not buy inventory efficiency here; it ran without fine-tuning or covariates.

6 DISCUSSION

The comparison metric changes what counts as a good forecast. With τ fixed at 0.90, the seven methods achieve M5 fill rates ranging from 0.935 to 0.967. Equal nominal quantiles are therefore not equal service outcomes. Matching realised fill rate allows stock requirements to be compared on the same basis. The VN1 window-to-window reversals also show why an inventory ranking based on a single test period may not hold in the next one [7].

On VN1, the retrospective demand split helps describe these reversals. TSB-NB uses less stock for stable-demand series, whereas LGB-Q reaches fill rate 0.92 on growing-demand series in two windows where TSB-NB cannot. Theodorou et al. [7] similarly find cases in which forecast accuracy and inventory cost favour different methods, though under different service and cost settings. In our experiment, the ranking also changes with the test period; no comparable reversal appears on M5 at the fill rates considered. If future demand growth were known in advance, assigning LGB-Q to growing series and TSB-NB to the rest would reduce inventory by 5.6–9.8% relative to LGB-Q on VN1. A classifier based on earlier features recovers only part of that hypothetical saving and remains unstable (repository).

Adaptive updating. The global models use the newest lags every week, but their parameters are re-estimated only every 13 weeks with fixed hyper-parameters, which can leave them one block behind a level shift on fast-growing VN1 series. More frequent or warm-started refits, heavier weights on recent observations, or re-tuning at each refit could reduce that lag; the VN1 sensitivity to tree size (Section 5.8) suggests configuration matters. We did not test these options, and neither the growth classifier nor learned combination weights (repository) beat simpler fixed choices, so added adaptivity needs out-of-sample validation.

How the quantiles are constructed also matters. On M5 at fill rate 0.94, LightGBM point forecasts with safety stock (LGB-T) or conformal adjustment (LGB-C) need 3.3–8.2% more inventory than directly estimated LightGBM quantiles. This is not directly comparable with the daily product–state experiment in [7], which constructed quantiles from residuals, or with the empirical-quantile results of [23]. Forecast construction and aggregation level are plausible sources of disagreement, but the studies do not isolate their effects.

Among the combinations, E7-med has the smallest maximum regret under the paper's descriptive metric (excess stock relative to the best single method, averaged over attainable fill-rate targets in each dataset–window run): 0.3% across six runs, versus 0.9% for the post hoc C2, 3.4% for E7, 9.4% for LGB-Q, 9.5% for LGB-T, 11.4% for HGB-Q and 20.7% for TSB-NB. LGB-T does not reach the highest target levels. When a lost sale costs at most 200 weeks of holding, C2's cost is within 3.1% of the cheapest single model in each run. At higher penalties the methods operate at $\tau = 0.99$; LGB-Q then has the highest fill rate, and C2 can cost as much as 29% more. These figures describe the six evaluated runs, not expected deployment performance.

Toward deployment. Because C2 is post hoc and E7-med was tested on only three windows per dataset, the next evaluation should be prospective: freeze the specification (members, equal weights, quantile grid, policy and metrics), apply it unchanged to later periods and to other retail datasets, and, in a company, run it in shadow mode beside the current method for several review cycles before any switch. Two guardrails follow from our results: use LGB-Q when lost sales cost more than about 200 weeks of holding, and monitor upper-quantile coverage, where TSB-NB under-covers.

Censored demand. Both datasets record sales, not demand, and contain no on-hand records. A sales-only detector that flags zero runs between two sales whose probability under the series' sale frequency p is below 1% ($(1 - p)^k < 0.01$) marks 7.4–9.0% of M5 and 6.1–6.9% of VN1 test weeks as suspected stockouts, about twice as often in intermittent and lumpy as in smooth series; because genuine pauses in demand are also flagged, these are upper bounds. Restricting the evaluation to the 72–82% of series without flagged weeks leaves the leading method of every window and C2's lowest maximum regret unchanged. A stricter scenario imputes each series' mean weekly sales in every flagged week. TSB-NB, whose occurrence probability falls during zero runs, then needs 17–27% more inventory than LGB-Q on M5 at fill rate 0.92; the leading method changes in three of six windows, always to a closely ranked method; and C2's maximum regret

rises to 9.8%, although it still needs 6.9–10.7% less inventory than LGB-Q on VN1. Censoring thus matters most for TSB-NB and for the post hoc case for C2 on brick-and-mortar data; stock-availability records would be needed to settle it. Censoring could also be handled in training rather than only in evaluation: flagged weeks could be treated as missing or right-censored observations and estimated with an expectation–maximisation step or a censored likelihood, as in lost-sales demand estimation [28], with censored quantile regression [29] for the quantile learners, or with a survival model of the time between sales. Because the flags are uncertain, such models would need on-hand records or a sensitivity analysis over the flagging threshold; we leave them to future work.

Limitations. Recorded sales stand in for demand, and stockouts are only inferred (see above). We did not tune hyperparameters; the M5 comparison persists across four LightGBM settings, but VN1 is sensitive to one of them (Section 5.8). Lead times of one, two and four weeks were examined in all three windows, without HGB-Q in the lead-time checks. Growth groups are formed retrospectively. Bootstrap intervals resample series rather than time periods, and each dataset supplies only three evaluation windows. Findings may differ for other retailers or inventory policies. Chronos-2 [22] was run zero-shot only; N-HiTS is an additional neural baseline (Section 5.8), not one of the seven methods in the main benchmark. Nabil et al. [27] forecast spare-part demand with a two-stage LightGBM/XGBoost model (an occurrence classifier times a size regressor) and set stock with a normal safety factor; a quantile version of this design is not tested here.

7 CONCLUSION

Under the weekly order-up-to policy tested here, LGB-Q requires the least inventory on M5 for a realised fill rate of 0.94 among the compared statistical and point-based alternatives. Across three test windows, its stock requirement is 3–11% lower, and it also has the lowest volume-weighted quantile loss. This finding persists across the reported lead-time checks and four LightGBM configurations.

The VN1 results do not identify a single leading method across periods. TSB-NB uses less stock on retrospectively identified stable-demand series, while LGB-Q achieves a 0.92 fill rate on growing series in two windows where TSB-NB cannot. Historical demand classes and the tested clustering do not produce consistent benefits when used to select a method prospectively. An equal-weight LGB-Q/TSB-NB quantile average reduces inventory by 6–11% relative to LGB-Q on VN1 at fill rate 0.92 and stays within 1.5% on M5; a pre-specified median of all seven methods saves 7–8% on VN1 with the lowest maximum regret, but no combination clearly beats LGB-Q on M5. Without an independent test period these findings are exploratory; the conclusions about single methods above do not depend on them.

Quantile liquidation reduces inventory on VN1 while changing fill rate by at most 0.02 percentage points. Its economic value is not established under the assumed costs over 26 weeks. Useful next tests would apply the quantile combination to new periods and retail datasets, select models using only information available before the forecast date, extend the liquidation horizon and compare more forecasting architectures under the same stock policy.

REFERENCES

[1] S. Makridakis, E. Spiliotis, and V. Assimakopoulos. 2022. The M5 competition: Background, organization, and implementation. International Journal of Forecasting 38, 4, 1325–1336. https://doi.org/10.1016/j.ijforecast.2021.07.007

[2] S. Makridakis, E. Spiliotis, and V. Assimakopoulos. 2022. M5 accuracy competition: Results, findings, and conclusions. International Journal of Forecasting 38, 4, 1346–1364. https://doi.org/10.1016/j.ijforecast.2021.11.013

[3] S. Makridakis, E. Spiliotis, V. Assimakopoulos, Z. Chen, A. Gaba, I. Tsetlin, and R. L. Winkler. 2022. The M5 uncertainty competition: Results, findings and conclusions. International Journal of Forecasting 38, 4, 1365–1385. https://doi.org/10.1016/j.ijforecast.2021.10.009

[4] M. Zanotti. 2025. The cost of ensembling: Is it always worth combining? arXiv:2506.04677.

[5] S. Wang, Y. Kang, E. Spiliotis, and F. Petropoulos. 2026. Multi-objective probabilistic forecast combination for inventory demand. arXiv:2606.04900.

[6] T. E. Goltsos, A. A. Syntetos, C. H. Glock, and G. Ioannou. 2022. Inventory–forecasting: Mind the gap. European Journal of Operational Research 299, 2, 397–419. https://doi.org/10.1016/j.ejor.2021.07.040

[7] E. Theodorou, E. Spiliotis, and V. Assimakopoulos. 2025. Forecast accuracy and inventory performance: Insights on their relationship from the M5 competition data. European Journal of Operational Research 322, 2, 414–426. https://doi.org/10.1016/j.ejor.2024.12.033

[8] Z. Mohammed, C. Anas, and M. El Hammoumi. 2026. A hybrid learning framework for forecasting uncertainty and adaptive inventory planning in retail supply chains. Supply Chain Analytics 13, 100180. https://doi.org/10.1016/j.sca.2025.100180

[9] O. Zabraoui, Y. Hmamou, A. Chafi, and S. Kammouri Alami. 2025. A comparative study of multi-algorithm optimization for inventory analytics in supply chains. Supply Chain Analytics 12, 100154. https://doi.org/10.1016/j.sca.2025.100154

[10] R. H. Teunter, A. A. Syntetos, and M. Z. Babai. 2011. Intermittent demand: Linking forecasting to inventory obsolescence. European Journal of Operational Research 214, 3, 606–615. https://doi.org/10.1016/j.ejor.2011.05.018

[11] A. A. Syntetos, J. E. Boylan, and J. D. Croston. 2005. On the categorization of demand patterns. Journal of the Operational Research Society 56, 5, 495–503. https://doi.org/10.1057/palgrave.jors.2601841

[12] S. Damato, N. Rubattu, D. Azzimonti, and G. Corani. 2026. Intermittent time series forecasting: Local vs global models. arXiv:2601.14031.

[13] N. Kourentzes, J. R. Trapero, and D. K. Barrow. 2020. Optimising forecasting models for inventory planning. International Journal of Production Economics 225, 107597. https://doi.org/10.1016/j.ijpe.2019.107597

[14] J. F. van der Haar, A. P. Wellens, R. N. Boute, and R. J. I. Basten. 2024. Supervised learning for integrated forecasting and inventory control. European Journal of Operational Research 319, 2, 573–586. https://doi.org/10.1016/j.ejor.2024.07.004

[15] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu. 2017. LightGBM: A highly efficient gradient boosting decision tree. In Advances in Neural Information Processing Systems 30.

[16] R. Koenker and G. Bassett. 1978. Regression quantiles. Econometrica 46, 1, 33–50. https://doi.org/10.2307/1913643

[17] J. Lei, M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman. 2018. Distribution-free predictive inference for regression. Journal of the American Statistical Association 113, 523, 1094–1111. https://doi.org/10.1080/01621459.2017.1307116

[18] R. Hyndman, A. Koehler, K. Ord, and R. Snyder. 2008. Forecasting with Exponential Smoothing. Springer. https://doi.org/10.1007/978-3-540-71918-2

[19] F. Pedregosa et al. 2011. Scikit-learn: Machine learning in Python. Journal of Machine Learning Research 12, 2825–2830.

[20] J. Demšar. 2006. Statistical comparisons of classifiers over multiple data sets. Journal of Machine Learning Research 7, 1–30.

[21] N. Vandeput. 2024. VN1 Forecasting – Accuracy Challenge. DataSource.ai.

[22] A. F. Ansari et al. 2025. Chronos-2: From univariate to universal forecasting. arXiv:2510.15821.

[23] E. Spiliotis, S. Makridakis, A. Kaltsounis, and V. Assimakopoulos. 2021. Product sales probabilistic forecasting: An empirical evaluation using the M5 competition data. International Journal of Production Economics 240, 108237. https://doi.org/10.1016/j.ijpe.2021.108237

[24] F. Petropoulos, X. Wang, and S. M. Disney. 2019. The inventory performance of forecasting methods: Evidence from the M3 competition data. International Journal of Forecasting 35, 251–265. https://doi.org/10.1016/j.ijforecast.2018.01.004

[25] M. Z. Babai, M. M. Ali, and K. Nikolopoulos. 2012. Impact of temporal aggregation on stock control performance of intermittent demand estimators: Empirical analysis. Omega 40, 713–721. https://doi.org/10.1016/j.omega.2011.09.004

[26] C. Challu, K. G. Olivares, B. N. Oreshkin, F. Garza, M. Mergenthaler-Canseco, and A. Dubrawski. 2023. N-HiTS: Neural hierarchical interpolation for time series forecasting. In Proceedings of the AAAI Conference on Artificial Intelligence 37, 6. https://doi.org/10.1609/aaai.v37i6.25854

[27] O. M. Nabil, N. H. Afia, and T. Ismail. 2026. Intermittent demand forecasting for spare parts using classification/regression models of machine learning and deep learning. Ain Shams Engineering Journal 17, 104445. https://doi.org/10.1016/j.asej.2026.104445

[28] S. Nahmias. 1994. Demand estimation in lost sales inventory systems. Naval Research Logistics 41, 6, 739–757. https://doi.org/10.1002/1520-6750(199410)41:6<739::AID-NAV3220410605>3.0.CO;2-A

[29] J. L. Powell. 1986. Censored regression quantiles. Journal of Econometrics 32, 1, 143–155. https://doi.org/10.1016/0304-4076(86)90016-3